

# Linux kernel

Ngoc Nguyen Tran

# Contents

- Introduction
- Features
- Linux kernel sources
- Building the kernel

# Linux kernel introduction

- The Linux kernel is one component of a system
- The Linux kernel was created as a hobby in 1991 by a Finnish student, Linus Torvalds

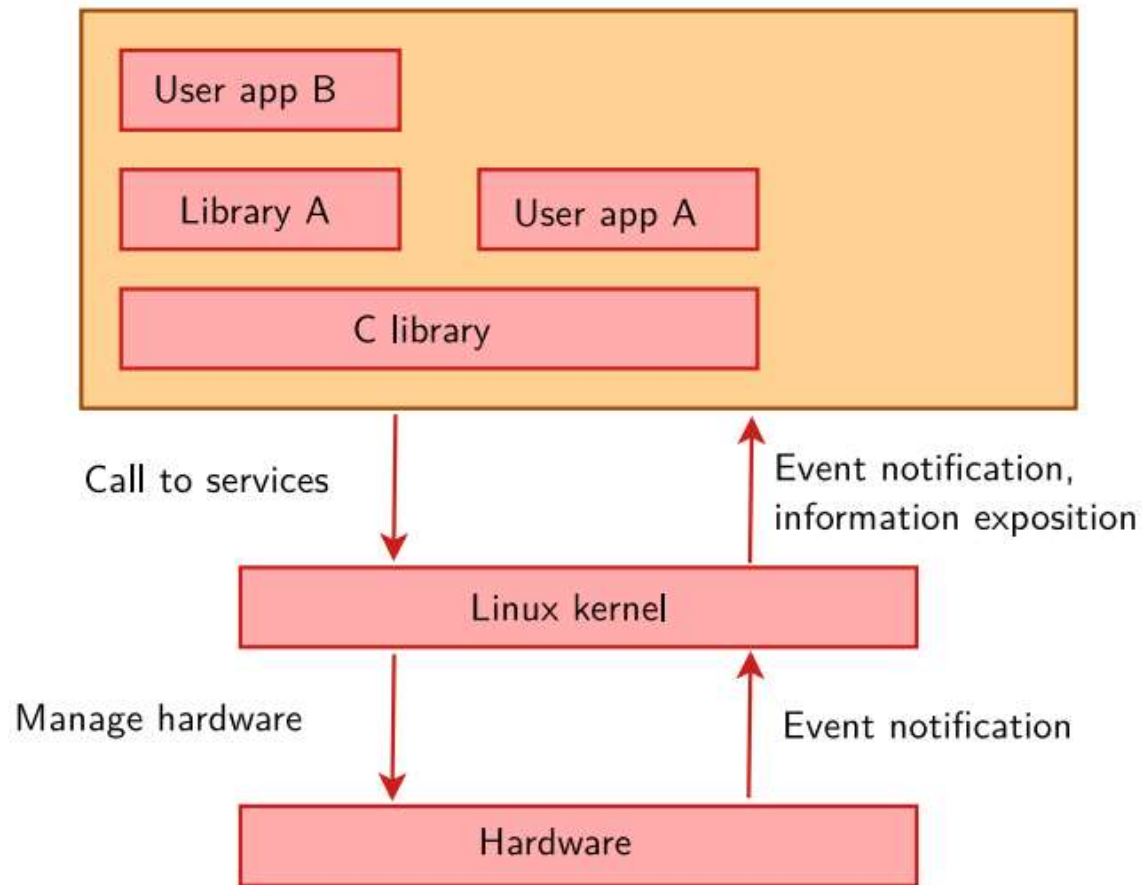


Linus Torvalds (2014)

# Linux kernel features

- Portability and hardware support.
  - Runs on most architectures.
- Scalability. Can run on super computers as well as on tiny devices (4 MB of RAM is enough).
- Compliance to standards and interoperability.
- Exhaustive networking support
- Security. It can't hide its flaws. Its
- code is reviewed by many experts.
- Stability and reliability.
- Modularity. Can include only what a system needs even at run time.
- Easy to program. You can learn from existing code. Many useful resources on the net.

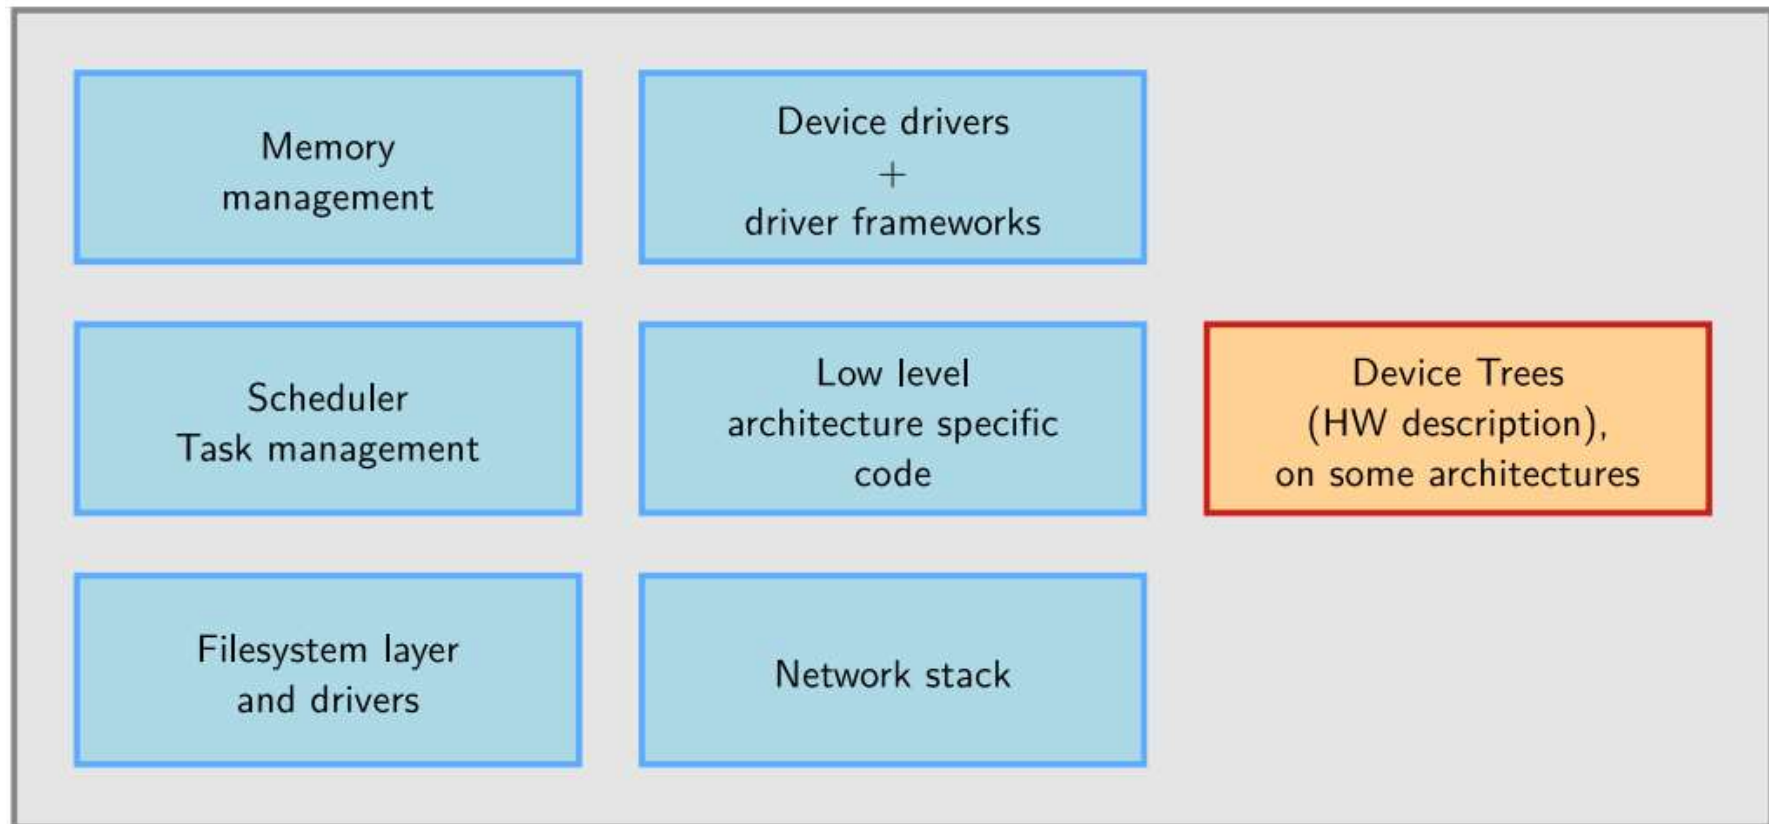
# Linux kernel in the system



# Pseudo filesystems

- Linux makes system and kernel information available in user space through **pseudo filesystems**, or **virtual filesystems**
- Pseudo filesystems allow applications to see directories and files that do not exist on any real storage: they are created and updated on the fly by the kernel
- The two most important pseudo filesystems are
  - proc, usually mounted on /proc:  
Operating system related information (processes, memory management parameters...)
  - sysfs, usually mounted on /sys:  
Representation of the system as a set of devices and buses. Information about these devices

# Inside the Linux kernel



# Linux kernel sources

- Available at <https://kernel.org>
- Many chip vendors supply their own kernel sources
- Many kernel sub-communities maintain their own kernel, with usually newer but less stable features
  - Architecture communities (ARM, MIPS, PowerPC, etc.), device drivers communities (I2C, SPI, USB, PCI, network, etc.), other communities (real-time, etc.)



# Getting Linux sources

- Available from <https://kernel.org/pub/linux/kernel>
  - Full tarballs: complete kernel sources -  
<https://kernel.org/pub/linux/kernel/v4.x/patch-4.20.xz>
  - Patches: differences between two kernel versions -  
<https://kernel.org/pub/linux/kernel/v4.x/linux-4.20.13.tar.xz>
- Get by git
- Using ***patch*** command to apply a patch to a current system
  - Need to run the patch command inside the toplevel kernel source directory

# Linux kernel size

- Linux 5.4 sources:
  - 66031 files (`git ls-files | wc -l`)
  - 27679764 lines (`git ls-files | xargs cat | wc -l`)
  - 889221135 bytes (`git ls-files | xargs cat | wc -c`)
- A minimum uncompressed Linux kernel just sizes 1-2 MB
- The Linux core (scheduler, memory management...) is pretty small!
- Why are these sources so big?

# Linux kernel size (2)

As of kernel version 4.6 (in percentage of number of lines).

- ▶ drivers/: 57.0%
- ▶ arch/: 16.3%
- ▶ fs/: 5.5%
- ▶ sound/: 4.4%
- ▶ net/: 4.3%
- ▶ include/: 3.5%
- ▶ Documentation/: 2.8%
- ▶ tools/: 1.3%
- ▶ kernel/: 1.2%
- ▶ firmware/: 0.6%
- ▶ lib/: 0.5%
- ▶ mm/: 0.5%
- ▶ scripts/: 0.4%
- ▶ crypto/: 0.4%
- ▶ security/: 0.3%
- ▶ block/: 0.1%
- ▶ ...

# Building the kernel

- The kernel configuration is the process of defining the set of options
  - Target architecture and device drivers
  - Your kernel capabilities (network, filesystem, real-time, etc.)
- Steps to build a kernel
  - Specifying the target architecture – under arch/:`export ARCH=arm`
  - Kernel configuration and build system – use `make`
  - Kernel configuration details - `.config` file at the root of kernel source
    - Edit manually or use tools such as **xconfig**, `gconfig`, `menuconfig`, `nconfig`
  - Initial configuration – find one configuration that works for your system first
    - Desktop: inside `/boot`
    - Embedded platform: default configuration are available for each CPU family, `/arch/<arch>/configs/`, the file `.config`
    - Run `make help` to find it
  - Create your own default configuration

# Kernel or module?

- The kernel image is a single file, resulting from the linking of all object files that correspond to features enabled in the configuration
  - This is the file that gets loaded in memory by the bootloader
  - All included features are therefore available as soon as the kernel starts, at a time where no filesystem exists
- Some features (device drivers, filesystems, etc.) can however be compiled as modules
  - These are plugins that can be loaded/unloaded dynamically to add/remove features to the kernel
  - Each module is stored as a separate file in the filesystem, and therefore access to a filesystem is mandatory to use modules
  - This is not possible in the early boot procedure of the kernel, because no filesystem is available

# Kernel option types

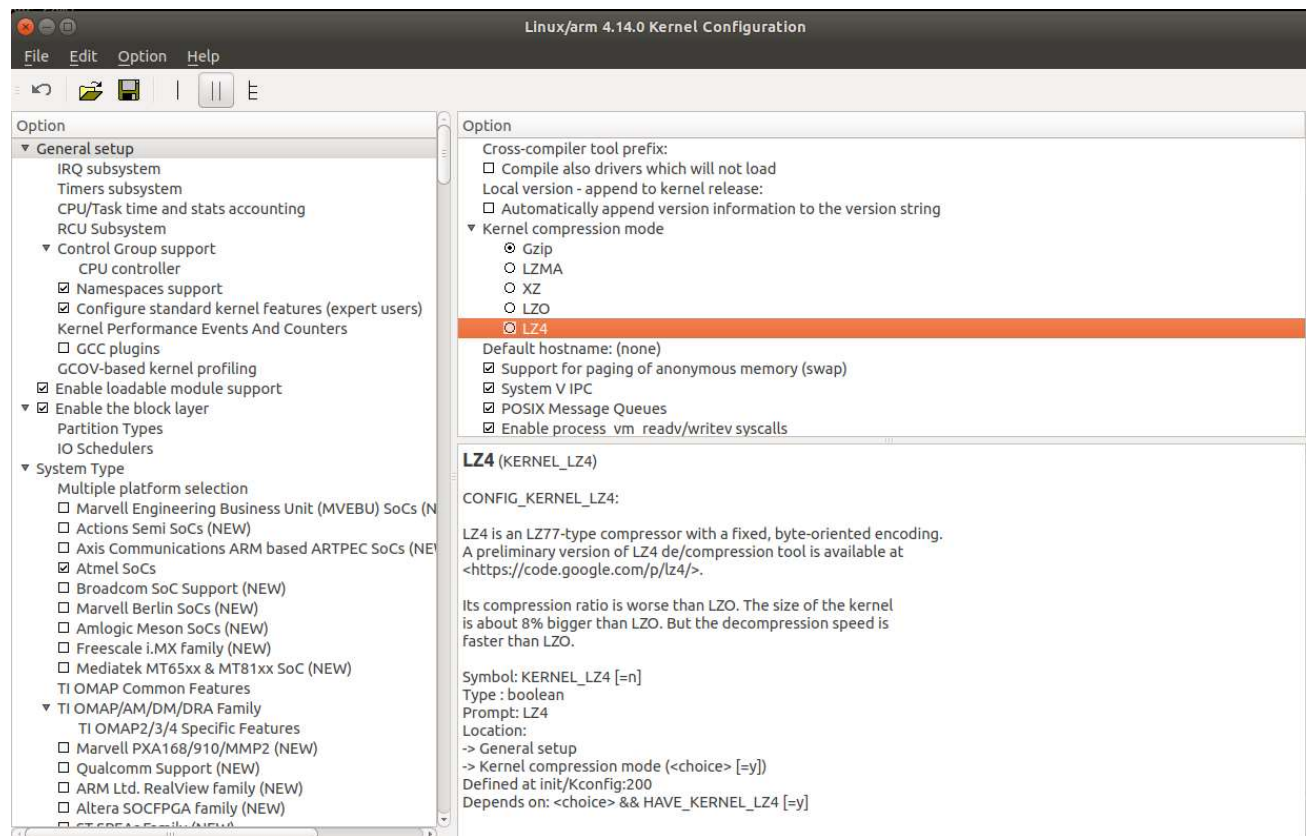
There are different types of options

- bool options, they are either
  - true (to include the feature in the kernel) or
  - false (to exclude the feature from the kernel)
- tristate options, they are either
  - true (to include the feature in the kernel image) or
  - module (to include the feature as a kernel module) or
  - false (to exclude the feature)
- int options, to specify integer values
- hex options, to specify hexadecimal values
- string options, to specify string values

# make xconfig

- make xconfig
  - The most common graphical interface to configure the kernel.
  - File browser: easier to load configuration files
  - Search interface to look for parameters
  - Required Debian / Ubuntu packages: qt5-default
- make gconfig
  - GTK based graphical configuration interface. Functionality similar to that of make xconfig.
  - Just lacking a search functionality.
  - Required Debian packages: libglade2-dev

# make xconfig screenshot





## Other similar tools

- gconfig
- menuconfig
- nconfig
- oldconfig

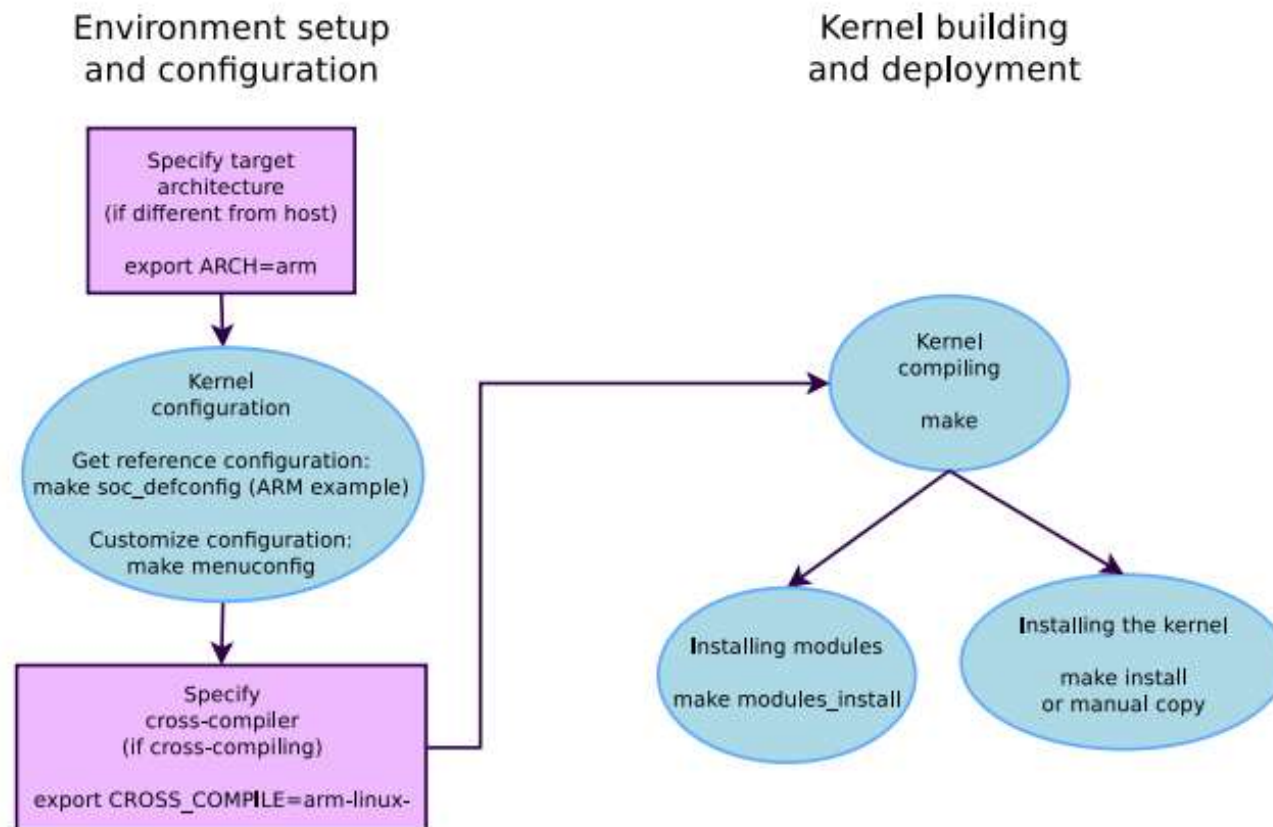
# Undoing configuration changes

- A frequent problem:
- After changing several kernel configuration settings, your kernel no longer works.
- If you don't remember all the changes you made, you can get back to your previous configuration:

```
$ cp .config.old .config
```

- All the configuration interfaces of the kernel (xconfig, menuconfig, oldconfig...) keep this .config.old backup copy.

# Kernel building overview



# Compiling a kernel

- Choose a compiler (*`$(CROSS_COMPILE)gcc`*)
  - Native: undefined `CROSS_COMPILE`
  - Cross-compiler → Example: *`export CROSS_COMPILE=arm-linux-gnueabi-`* for *`arm-linux-gnueabi-gcc`*
- Specifying ARCH and CROSS\_COMPILE
  - Pass *ARCH* and *CROSS\_COMPILE* on the make command line:  
Example: *`make ARCH=arm CROSS_COMPILE=arm-linux-`*
  - Define *ARCH* and *CROSS\_COMPILE* as environment variables: (*export*)
- Kernel compilation
  - `make`
    - Can run multiple jobs in parallel if having multiple CPU cores (`make -j 8`)
  - Generates:
    - `vmlinux`, the raw uncompressed kernel image (cannot be booted)
    - `arch/<arch>/boot/*Image`, the final, usually compressed, kernel image that can be booted  
Example: `bzImage` for x86, `zImage` for ARM, `vmlinux.bin.gz` for ARC, etc.

# Kernel installation – native case

- make install
  - Does the installation for the host system by default, so needs to be run as root.
- Installs
  - */boot/vmlinuz-<version>*
    - Compressed kernel image. Same as the one in *arch/<arch>/boot*
  - */boot/System.map-<version>*
    - Stores kernel symbol addresses for debugging purposes (obsolete: such information is usually stored in the kernel itself)
  - */boot/config-<version>*
    - Kernel configuration for this version
- In GNU/Linux distributions, typically re-runs the bootloader configuration utility to make the new kernel available at the next boot.

# Kernel installation: embedded case

- make install is rarely used in embedded development, as the kernel image is a single file, easy to handle.
- Another reason is that there is no standard way to deploy and use the kernel image.
- Therefore making the kernel image available to the target is usually manual or done through scripts in build systems.
- It is however possible to customize the make install behavior in  
arch/<arch>/boot/install.sh

# Kernel cleanup target

- Clean-up generated files (to force re-compilation):  
    `make clean`
- Remove all generated files. Needed when switching from one architecture to another. Caution: it also removes your `.config` file!  
    `make mrproper`
- Also remove editor backup and patch reject files (mainly to generate patches):  
    `make distclean`
- If you are in a git tree, remove all files not tracked (and ignored) by git:  
    `git clean -fdx`